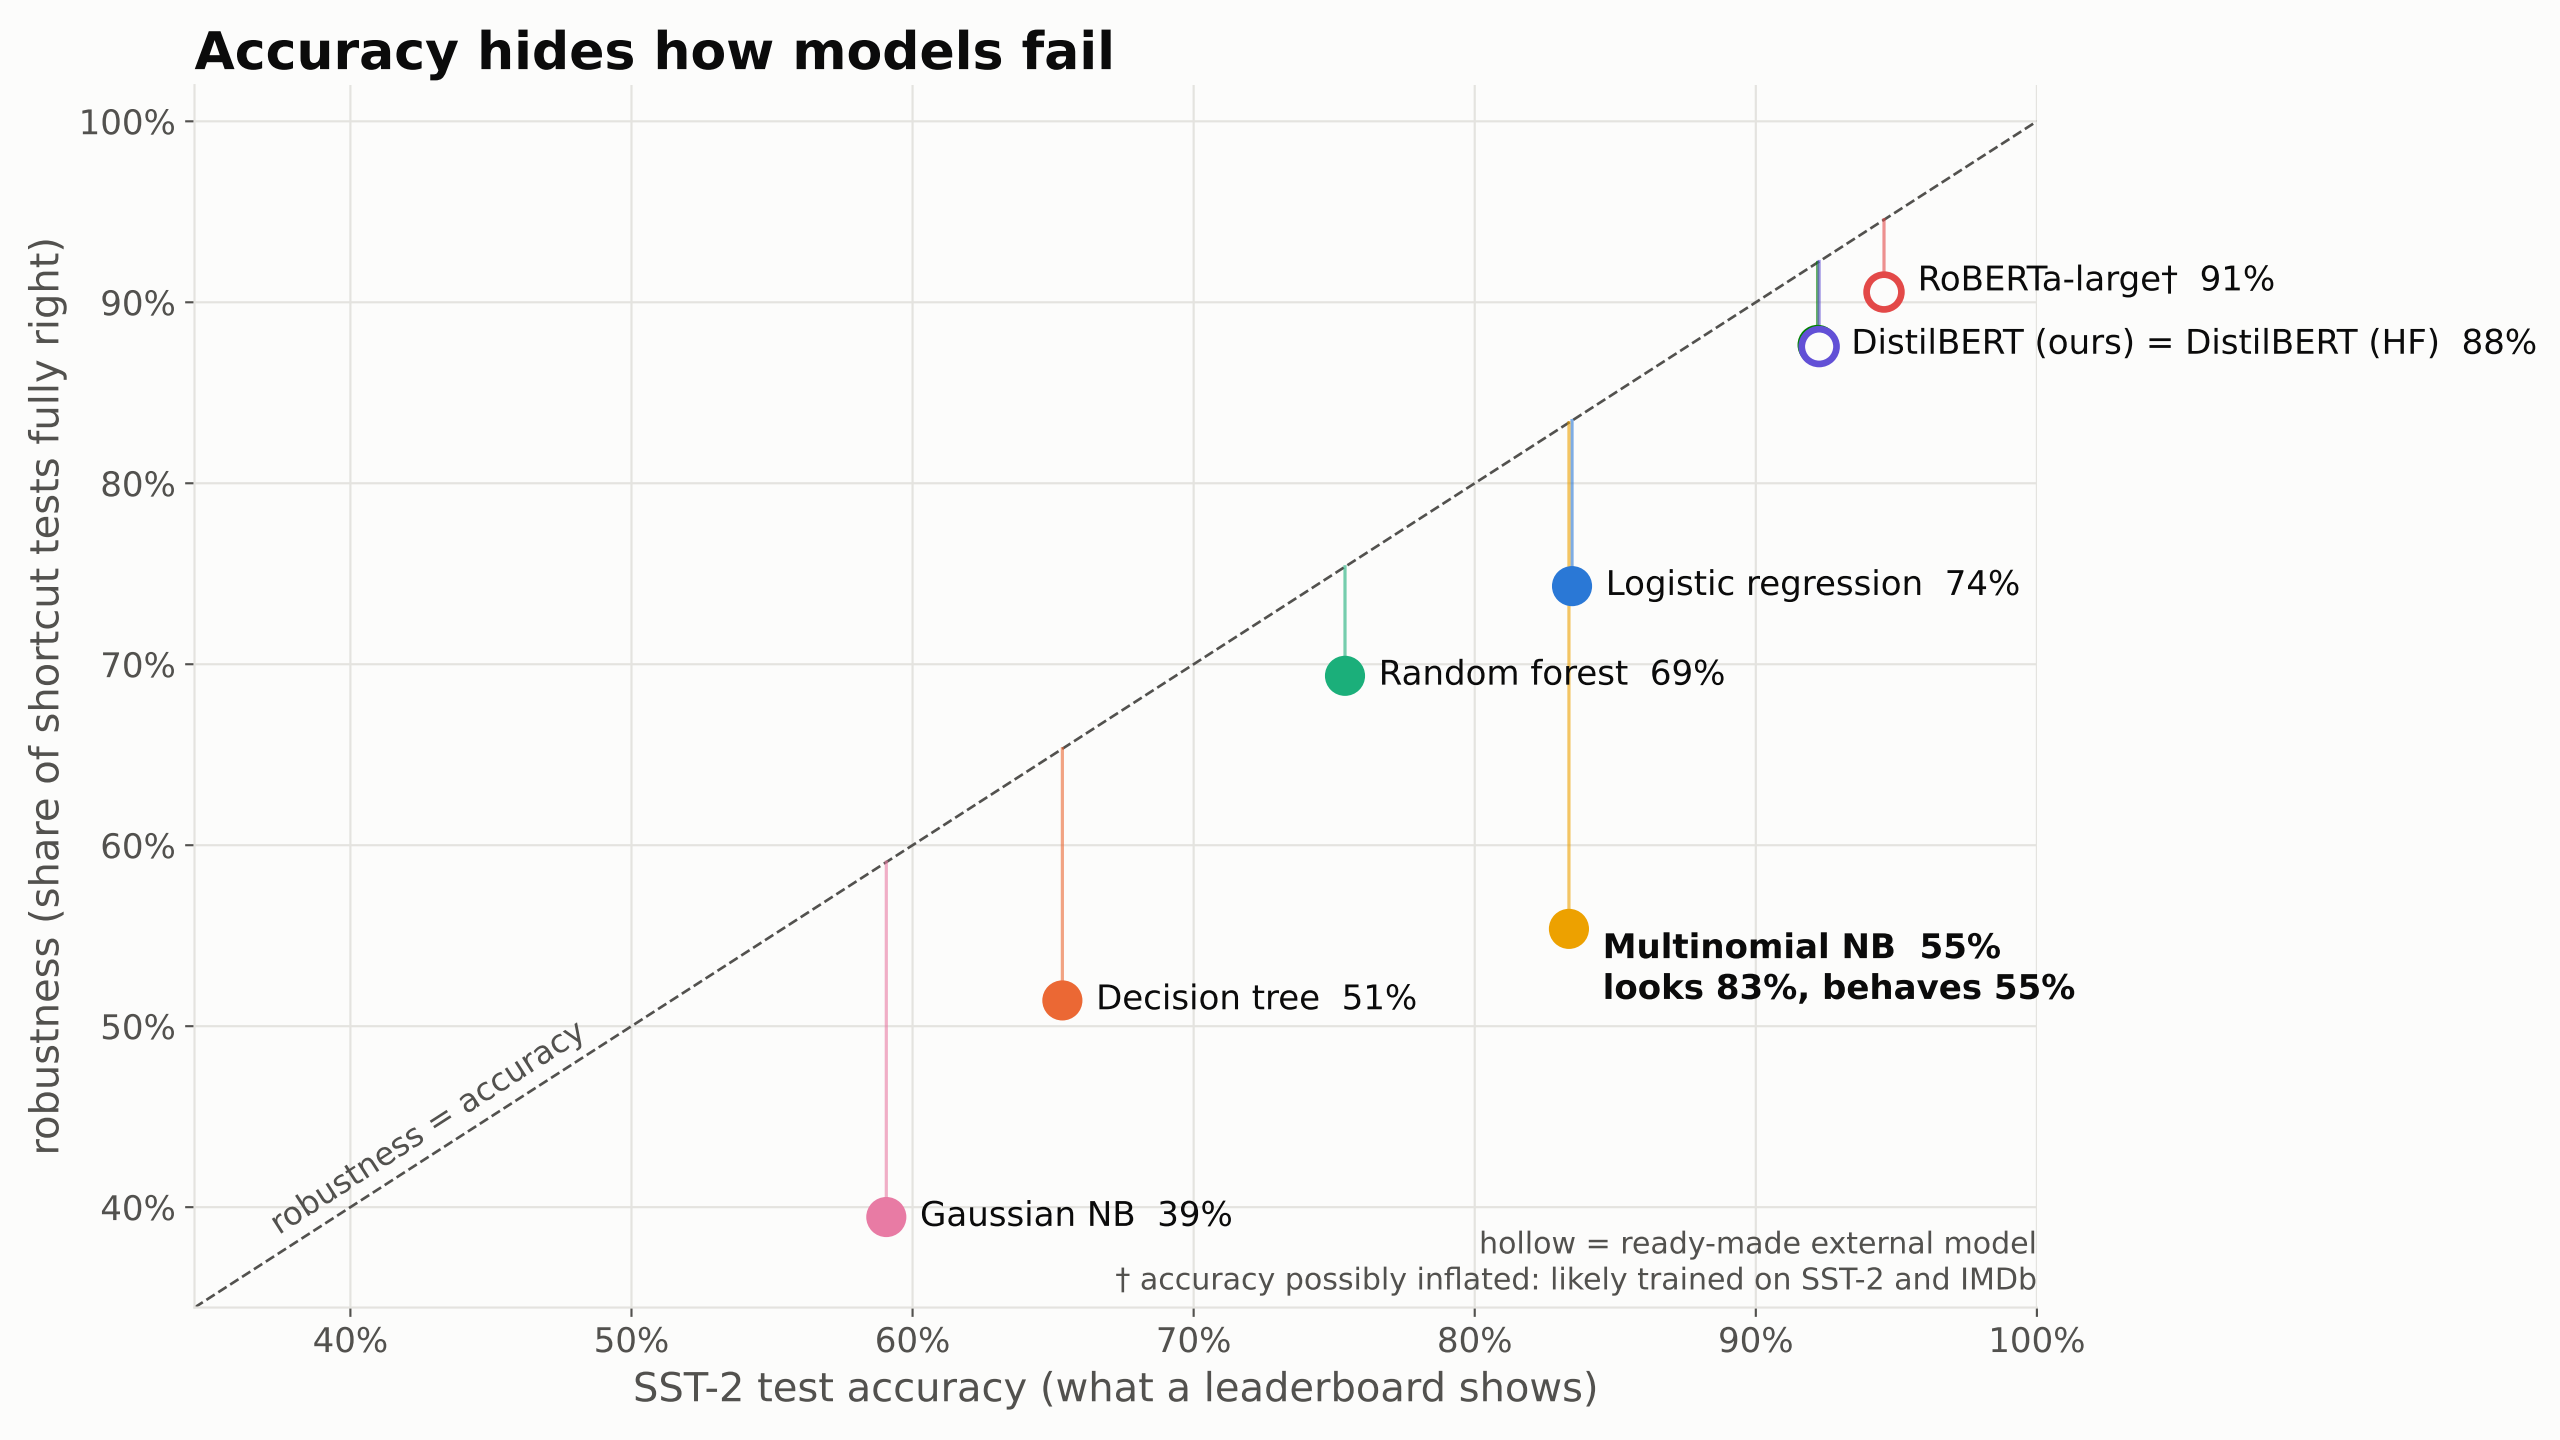
Accuracy hides how models fail
robustness (share of shortcut tests fully right)
SST-2 test accuracy (what a leaderboard shows)
100%
90%
80%
70%
60%
50%
40%
40%
50%
60%
70%
80%
90%
100%
robustness = accuracy
RoBERTa-large† 91%
DistilBERT (ours) = DistilBERT (HF) 88%
Logistic regression 74%
Random forest 69%
Multinomial NB 55%
looks 83%, behaves 55%
Decision tree 51%
Gaussian NB 39%
hollow = ready-made external model
† accuracy possibly inflated: likely trained on SST-2 and IMDb